

# STAT S4201 001, Homework 2

Brian Weinstein (bmw2148)

Feb 10, 2016

Code is attached here and also posted at <https://github.com/BrianWeinstein/advanced-data-analysis>. Where relevant, code snippets and output are included in-line.

## Problem 1:

- (a) Suppose that  $\{X_1, X_2, \dots, X_n\}$  is a random sample from  $N(\mu, \sigma^2)$ . Construct a 95% confidence interval for  $\sigma^2$  under the following scenarios: (a)  $\mu$  is known to be 0. (b)  $\mu$  is unknown.

Since the  $X_i$  are normal, the construction

$$\frac{(n-1)s^2}{\sigma^2} \sim \chi_{n-1}^2$$

has a Chi-Squared distribution with  $(n-1)$  degrees of freedom, where  $s^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$  and  $\sigma^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2$ .

Therefore, we can construct a  $100(1 - \alpha)\%$  confidence interval for  $\sigma^2$  as

$$\begin{aligned} 1 - \alpha &= P \left[ F_{n-1} \left( \frac{\alpha}{2} \right) \leq \frac{(n-1)s^2}{\sigma^2} \leq F_{n-1} \left( 1 - \frac{\alpha}{2} \right) \right] \\ &= P \left[ \frac{1}{F_{n-1} \left( \frac{\alpha}{2} \right)} \geq \frac{\sigma^2}{(n-1)s^2} \geq \frac{1}{F_{n-1} \left( 1 - \frac{\alpha}{2} \right)} \right] \\ &= P \left[ \frac{(n-1)s^2}{F_{n-1} \left( 1 - \frac{\alpha}{2} \right)} \leq \sigma^2 \leq \frac{(n-1)s^2}{F_{n-1} \left( \frac{\alpha}{2} \right)} \right] \end{aligned}$$

where  $F_{n-1}(x)$  is the cumulative distribution function of  $\chi_{n-1}^2$ .

Setting  $\alpha = 0.05$  to construct a 95% CI, we have

$$0.95 = P \left[ \frac{(n-1)s^2}{F_{n-1}(0.975)} \leq \sigma^2 \leq \frac{(n-1)s^2}{F_{n-1}(0.025)} \right].$$

- (b) Fix  $n = 10$  and  $\sigma = 1$ . Run a Monte Carlo simulation to confirm that the confidence interval you constructed under the scenario (a) produces a coverage of 95%. Report how many random samples were drawn in your simulation and how close your coverage was to 95%.

See attached code. In my simulation I drew 1000 random samples, with a 95% confidence interval capturing the true variance in 948 of the trials (94.8% coverage).

**Problem 2: Ramsey 3.22**

```
# Data input
time26 <- c(5.79, 1579.52, 2323.70)
time28 <- c(68.8, 108.29, 110.29, 426.07, 1067.60)
```

(a) Form two new variables by taking the logarithms of the breakdown times.

```
> Y1 <- log(time26) ; Y1
[1] 1.756132 7.364876 7.750916
> Y2 <- log(time28) ; Y2
[1] 4.231204 4.684813 4.703113 6.054604 6.973168
```

(b)  $\bar{Y}_1 - \bar{Y}_2 = 5.6240 - 5.3294 = 0.2946$

(c)  $\exp(\bar{Y}_1 - \bar{Y}_2) = \exp(0.2946) = 1.3426$ , where 1.3426 is the multiplicative treatment effect, indicating that the breakdown time at 26 kV is estimated to be 1.3426 times larger than the breakdown time at 28 kV.

(d) Compute a 95% confidence interval for the difference in mean log breakdown times. Take the antilogarithms of the endpoints and express the result in a sentence.

The pooled standard deviation is given by

$$s_p = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{(n_1 + n_2 - 2)}} = \sqrt{\frac{(3 - 1)(11.257) + (5 - 1)(1.310)}{(3 + 5 - 2)}} = 2.1508,$$

where  $n_1 = 3$  and  $n_2 = 5$  are the sample sizes and  $s_1^2 = 11.2574$  and  $s_2^2 = 1.3104$  are the sample variances of the log-transformed measurements.

The standard error of  $(\bar{Y}_1 - \bar{Y}_2)$  is given by

$$SE(\bar{Y}_1 - \bar{Y}_2) = s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}} = (2.1508) \sqrt{\frac{1}{3} + \frac{1}{5}} = 1.5707.$$

Therefore a 95% confidence interval for the difference in mean log breakdown times is

$$\begin{aligned} & (\bar{Y}_1 - \bar{Y}_2) \pm t_{(3+5-2)} \left(1 - \frac{\alpha}{2}\right) SE(\bar{Y}_1 - \bar{Y}_2) \\ & 0.2946 \pm t_6 \left(1 - \frac{0.05}{2}\right) 1.5707 \\ & 0.2946 \pm (2.4469)(1.5707) \\ & 0.2946 \pm 3.8435 \\ & \implies -3.5489 \leq (\bar{Y}_1 - \bar{Y}_2) \leq 4.1381. \end{aligned}$$

Taking the anilogarithms of the confidence interval endpoints,

$$\text{Lower confidence limit} = e^{-3.5489} = 0.0288$$

$$\text{Upper confidence limit} = e^{4.1381} = 62.682,$$

we find that the breakdown time at 26 kV is estimated to be 1.3426 (from part c) times longer than the breakdown time at 28 kV, (95% confidence: 0.0288 to 62.682 times).

**Problem 3: Ramsey 3.25**

See attached code. When excluding either (1) no observations, (2) observation 646, or (3) observations 646 and 645, there is no evidence that the mean dioxin level in Vietnam veterans is greater than the mean dioxin level in non-Vietnam veterans. The difference in means (non-veteran minus veteran) (parts per trillion), one-sided p-values, and 95% confidence intervals for the difference in means (parts per trillion) are:

Case	Difference in means	one-sided p-value	95% CI
(1)	-0.0745	0.3963	-0.6305 to 0.4815
(2)	-0.0113	0.4805	-0.6305 to 0.4815
(3)	0.0210	0.5386	-0.6305 to 0.4815

**Problem 4: Ramsey 3.28**

Boxplots show in Figure 1. When including all of the observations, the t-test generates a two-sided p-value of 0.0809:

```
> t.test(formula=Humerus~Status, data=sparrowData,
+ var.equal=TRUE, conf.level=0.95)

Two Sample t-test

data:  Humerus by Status
t = -1.777, df = 57, p-value = 0.0809
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.021446053  0.001279386
sample estimates:
mean in group Perished mean in group Survived
      0.7279167          0.7380000
```

And when excluding the smallest length in the Perished group, we have a two-sided p-value of 0.18:

```
> t.test(formula=Humerus~Status, data=sparrowData,
+ subset=!(Humerus==min(sparrowData$Humerus) & Status=="Perished"),
+ var.equal=TRUE, conf.level=0.95)

Two Sample t-test

data:  Humerus by Status
t = -1.3578, df = 56, p-value = 0.18
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.017542698  0.003368785
sample estimates:
mean in group Perished mean in group Survived
      0.730913          0.738000
```

With the full dataset, the p-value indicates that the data is suggestive, but isn't conclusive, in the hypothesis that the humerus length differs in the two populations. When excluding the smallest length in the Perished group, the p-value indicates that there's little evidence of a difference between the groups.

Since the conclusion of the study depends on which dataset is used, we must either use resistant analysis (see Problem 6) or report the results of both analyses, as done here.




Figure 1: Humerus lengths of male sparrows that perished or survived in a winter storm. The smallest length in the Perished group (indicated with a dot) is an extreme outlier and is later excluded from the analysis.




Figure 2: Out-of-state vs in-state tuition costs for 25 public colleges.

### Problem 5: Ramsey 3.32

Analyze the data to describe the extent to which:

- (a) *public school tuition is more expensive when out-of-state than when in-state.*

Boxplots of out-of-state vs in-state tuition for the 25 public schools in Figure 2 reveal that (1) out-of-state tuition has a higher center than in-state tuition, (2) out-of-state tuition has a higher standard deviation than in-state tuition, and (3) some outliers may be skewing the results.

The standard error for out-of-state tuition is 4,500.41 USD, which is 1.71 times higher than the standard error for in-state tuition of 2633.68 USD, making the two-sample t-tools unreliable here.

We could try log-transforming the data before using the two-sample t-tools, but this isn't necessary here. Instead, the one-sample paired t-tools may be more appropriate in this case, since each {out-of-state, in-state} measurement is an observation on a single college. The distribution of the the difference between each {out-of-state, in-state} pair is




Figure 3: (Out-of-state) - (in-state) tuition costs for 25 public colleges.




Figure 4: In-state tuition costs for private vs public colleges.

nearly normal, as shown in Figure 3.

Using the paired t-test to test the alternative hypothesis that (out-of-state tuition) – (in-state tuition)  $> 0$ , we find substantial evidence that the mean difference in the out-of-state tuition and in-state tuition among a random sample of 25 public colleges is positive (one-sided p-value =  $3.264 \times 10^{-13}$ ). It's estimated that the mean tuition cost is \$9,038.4 more for out-of-state tuition (95% confidence interval: \$7,687.08 to \$10,389.72).

(b) *in-state tuition is more expensive for private schools than for public schools.*

Boxplots of in-state tuition costs for private vs public colleges in Figure 4 reveal that private tuition has both a higher center and significantly higher spread than public tuition, making this dataset a good candidate for a log transformation.

The log-transformed data is shown in Figure 5, where the private tuition has a standard error of 0.495 and the public tuition has a similar standard error of 0.334 (on the log scale). Since the sample sizes in each group are the same, the the difference in standard errors has little effect on the robustness of the t-tools.

Using the two-sample t-test on the log-transformed tuitions, we find substantial evidence that private school in-state tuition is greater than public school in-state tuition on the log scale (one-sided p-value:  $6.31 \times 10^{-15}$ ). Back-transforming the estimate and confidence




Figure 5: Log-transformed in-state tuition costs for private vs public colleges.

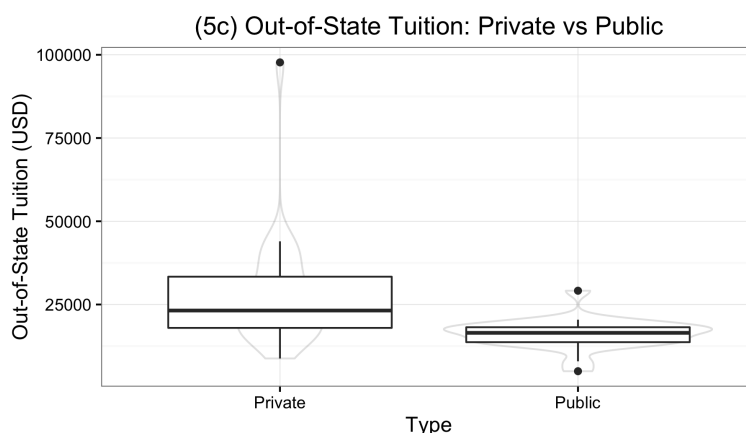


Figure 6: In-state tuition costs for private vs public colleges.

interval from the log scale to the original scale, it's estimated that the mean tuition cost is 3.692 times higher for out-of-state than in-state (95% confidence interval: 2.904 to 4.695 times).

- (c) *out-of-state tuition is more expensive for private schools than for public schools.*

Boxplots of out-of-state tuition costs for private vs public colleges in Figure 6 reveal that private tuition has both a higher center and significantly higher spread than public tuition, making this dataset a good candidate for a log transformation.

The log-transformed data is shown in Figure 7, where the private tuition has a standard error of 0.495 and the public tuition has a similar standard error of 0.334 (on the log scale). Since the sample sizes in each group are the same, the difference in standard errors has little effect on the robustness of the t-tools.

Using the two-sample t-test on the log-transformed tuitions, we find substantial evidence that private school in-state tuition is greater than public school in-state tuition on the log scale (one-sided p-value: 0.0001073). Back-transforming the estimate and confidence interval from the log scale to the original scale, it's estimated that the mean tuition cost is 1.613 times higher for out-of-state than in-state (95% confidence interval: 1.269 to 2.0512 times).




Figure 7: Log-transformed out-of-state tuition costs for private vs public colleges.

### Problem 6: Ramsey 4.19

Using the `wilcox.test` R function:

```
> # load data
> sparrowData <- Sleuth3::ex0221
>
> # perform rank-sum test
> wilcox.test(formula=Humerus~Status, data=sparrowData,
+             paired=FALSE, # perform rank-sum test
+             correct=TRUE, # apply continuity correction in the normal approx for the p-value
+             conf.int=TRUE, conf.level=0.95)

Wilcoxon rank sum test with continuity correction

data:  Humerus by Status
W = 331, p-value = 0.1718
alternative hypothesis: true location shift is not equal to 0
95 percent confidence interval:
 -0.019076357  0.003033795
sample estimates:
difference in location
 -0.007022596

Warning messages:
1: In wilcox.test.default(x = c(0.659, 0.689, 0.703, 0.702, 0.709, :
  cannot compute exact p-value with ties
2: In wilcox.test.default(x = c(0.659, 0.689, 0.703, 0.702, 0.709, :
  cannot compute exact confidence intervals with ties
```

- The two-sided p-value is 0.1718.
- From the documentation, by default, when the `exact` argument isn't specified the `wilcox.test` function computes an exact p-value only when the combined sample has fewer than 50 values and there are no ties. Otherwise, as is the case here, since our dataset has 59 observations and also has ties, a normal approximation is used.
- The p-value based on the normal approximation uses a continuity correction to the Z-statistic.
- The (two-sided) p-value from the rank-sum test (0.1718) is greater than the p-value from

the two-sample t-test (0.0809) and slightly smaller than the p-value from the two-sample t-test when the smallest observation is excluded (0.18) in Problem 4.

- (e) In this case, the (two-sided) p-value obtained from the two-sample t-test when excluding outliers is nearly identical to the p-value obtained in the rank-sum test.

The two-sample t-test assumes nearly normal populations, equal population variances, and independence of samples. All of these assumptions are nearly satisfied both when the outlier is and isn't excluded.

By design, the rank-sum test is less efficient when the populations are normal, but is much better at dealing with outliers and makes no assumption about the population distributions.

Removing the outlier isn't entirely justified here — there isn't any indication that there were recording errors or that the outlier came from a different population — so the rank-sum test may be a better choice, since it generates a similar p-value and entirely avoids the issue of excluding outliers.



```
#####

# Brian Weinstein - bmw2148
# STAT S4201 001
# Homework 2
# 2016-02-10

# set working directory
setwd("~/Documents/advanced-data-analysis/homework_02")

# load packages
library(dplyr)
library(Sleuth3) # Data sets from Ramsey and Schafer's "Statistical Sleuth
(3rd ed)"
library(ggplot2); theme_set(theme_bw())
library(tidyr)

# Problem 1
#####

# Define a fn to calculate the conf interval for the sample var, given the
sample var, n, and conf level
MakeVarConfidenceInterval <- function(sampleVar, n, confLevel){

  sigLevel <- 1 - confLevel

  lowerBound <- (n-1) * sampleVar / qchisq(p=(1-(sigLevel/2)), df=(n-1),
lower.tail=TRUE)
  upperBound <- (n-1) * sampleVar / qchisq(p=(sigLevel/2), df=(n-1),
lower.tail=TRUE)

  confidenceInterval <- c(lowerBound, upperBound)

  return(confidenceInterval)

}

# define variables for this problem
variance <- 1
mean <- 0
sample_size <- 10

# Initialize a counter at 0
numTrialsCapturingTrueVar=0

# Run 1000 trials, checking if the 95% CI captures the true varaince
set.seed(1)
for(trial in 1:1000){

  # Generate a random sample from N(mu=0, var=1)
  random_sample <- rnorm(n=sample_size, mean=mean, sd=sqrt(variance))

```

```

# Compute the sample var of the random sample
sample_var <- var(random_sample)

# Compute a 95% CI for the sample variance
confInt <- MakeVarConfidenceInterval(sampleVar=sample_var, n=sample_size,
confLevel=0.95)

# Check if CI captured true sample variance and increment counter
numTrialsCapturingTrueVar = numTrialsCapturingTrueVar + ifelse(confInt[1] <
variance & confInt[2] > variance, 1, 0)

}

c(numTrialsCapturingTrueVar, trial)
numTrialsCapturingTrueVar/trial

rm(list = ls()) # clear working environment

# Problem 2: Ramsey 3.22
#####

# Data input
time26 <- c(5.79, 1579.52, 2323.70)
time28 <- c(68.8, 108.29, 110.29, 426.07, 1067.60)

# Part a ### ### ### ### ### ### ### ### ### ### ### ### ### ### ### ###
###

Y1 <- log(time26) ; Y1
Y2 <- log(time28) ; Y2

# Part b ### ### ### ### ### ### ### ### ### ### ### ### ### ### ### ###
###

mean(Y1)
mean(Y2)

mean(Y1) - mean(Y2)

# Part c ### ### ### ### ### ### ### ### ### ### ### ### ### ### ### ###
###

exp(mean(Y1) - mean(Y2))

# Part d ### ### ### ### ### ### ### ### ### ### ### ### ### ### ### ###
###

# Define variables to calculate confidence intervals
logDiff <- mean(Y1) - mean(Y2)
pooledSD <- sqrt( ((length(Y1) - 1)*var(Y1) + (length(Y2) - 1)*var(Y2)) /
(length(Y1) + length(Y2) - 2) )
se <- pooledSD * sqrt((1/3) + (1/5))

```

```

quantile <- qt(1-(0.05/2), df=(3+5-2))

# Compute confidence interval
logDiff - (quantile * se)
logDiff + (quantile * se)

# Compute anilogarithm of confidence interval
exp(logDiff - (quantile * se))
exp(logDiff + (quantile * se))

rm(list = ls()) # clear working environment

# Problem 3: Ramsey 3.25
#####

# load data
agentOrangeData <- Sleuth3::case0302

# compute group difference, 1-sided p-value, and 95% CI; with all observations
t1 <- t.test(formula=Dioxin~Veteran, data=agentOrangeData,
              var.equal=TRUE, conf.level=0.95, alternative="less")
-diff(t1$estimate)[[1]]
t1$p.value
t.test(formula=Dioxin~Veteran, data=agentOrangeData,
        var.equal=TRUE, conf.level=0.95)$conf.int

# compute group difference, 1-sided p-value, and 95% CI; excluding observation
646
t2 <- t.test(formula=Dioxin~Veteran, data=agentOrangeData[-646, ],
              var.equal=TRUE, conf.level=0.95, alternative="less")
-diff(t2$estimate)[[1]]
t2$p.value
t.test(formula=Dioxin~Veteran, data=agentOrangeData,
        var.equal=TRUE, conf.level=0.95)$conf.int

# compute group difference, 1-sided p-value, and 95% CI; excluding
observations 646 and 645
t3 <- t.test(formula=Dioxin~Veteran, data=agentOrangeData[-(645:646), ],
              var.equal=TRUE, conf.level=0.95, alternative="less")
-diff(t3$estimate)[[1]]
t3$p.value
t.test(formula=Dioxin~Veteran, data=agentOrangeData,
        var.equal=TRUE, conf.level=0.95)$conf.int

rm(list = ls()) # clear working environment

# Problem 4: Ramsey 3.28
#####

# load data
sparrowData <- Sleuth3::ex0221

```

```

# create boxplots
ggplot(sparrowData, aes(x=Status, y=Humerus)) +
  geom_violin(alpha=0.15) +
  geom_boxplot() +
  labs(y="Humerus Length (inches)", title="Humerus Length: Perished vs
Survived")
ggsave(filename="writeup/4_full.png", width=6.125, height=3.5, units="in")

# compute group difference, 2-sided p-value, and 95% CI; with all observations
t.test(formula=Humerus~Status, data=sparrowData,
       var.equal=TRUE, conf.level=0.95)

# compute group difference, 2-sided p-value, and 95% CI; excluding the
smallest length in the perished group
t.test(formula=Humerus~Status, data=sparrowData,
       subset=!(Humerus==min(sparrowData$Humerus) & Status=="Perished"),
       var.equal=TRUE, conf.level=0.95)

rm(list = ls()) # clear working environment

# Problem 5: Ramsey 3.32
#####

# load data
tuitionData <- Sleuth3::ex0332

# Part a: Public school tuition (out-of-state > in-state) ### ### ### ### ###
### ### ### ### ### ### ### ### ### ### ### ### ### ### ###

# Subset the data for part a and convert to long format
tuitionData.a <- tuitionData %>%
  filter(Type=="Public") %>%
  gather(key="tuitionType", value="tuition", c(OutOfState, InState),
factor_key=TRUE) %>%
  arrange(College, tuitionType)

# create boxplots
ggplot(tuitionData.a, aes(x=tuitionType, y=tuition)) +
  geom_violin(alpha=0.15) +
  geom_boxplot() +
  labs(y="Tuition (USD)", title="(5a) Public School Tuition: In-State vs Out-
of-State")
ggsave(filename="writeup/5a_twoSample.png", width=6.125, height=3.5,
units="in")

# check group standard errors
tuitionData.a %>%
  group_by(tuitionType) %>%
  summarize(sd=sd(tuition))

# group standard errors are significantly different, so two-sample t-test is
unreliable here

```

```

# since the measurements are each paired, we can try the one-sample paired t-
test

# Create a new dataframe with the tuition differences
tuitionData.a2 <- tuitionData %>%
  filter(Type=="Public") %>%
  mutate(diff=OutOfState-InState)

# create boxplot
ggplot(tuitionData.a2, aes(x=Type, y=diff)) +
  geom_violin(alpha=0.15) +
  geom_boxplot() +
  labs(y="Tuition (USD)", title="(5a) Public School Tuition: (Out-of-State) -
(In-State)")
ggsave(filename="writeup/5a_paired.png", width=6.125, height=3.5, units="in")

# Perform paired t-test
t.test(x=tuitionData.a2$OutOfState, y=tuitionData.a2$InState, paired=TRUE,
  conf.level=0.95, alternative="greater")
# Perform a two sided paired t-test for the confidence interval
t.test(x=tuitionData.a2$OutOfState, y=tuitionData.a2$InState, paired=TRUE,
  conf.level=0.95, alternative="two.sided")$conf.int

# Part b: In-state tutition (private > public) #### ### ### ### ### ### ### ###
### ### ### ### ### ### ### ### ### ### ### ###

# Create dataset for part b
tuitionData.b <- tuitionData %>%
  mutate(OutOfState=NULL)

# create boxplots
ggplot(tuitionData.b, aes(x=Type, y=InState)) +
  geom_violin(alpha=0.15) +
  geom_boxplot() +
  labs(y="In-State Tuition (USD)", title="(5b) In-State Tuition: Private vs
Public")
ggsave(filename="writeup/5b_original.png", width=6.125, height=3.5,
units="in")

# Create a log-transformed InState tuition column
tuitionData.b <- tuitionData.b %>%
  mutate(LogInState=log(InState))

# create boxplots of the log-transformed data
ggplot(tuitionData.b, aes(x=Type, y=LogInState)) +
  geom_violin(alpha=0.15) +
  geom_boxplot() +
  labs(y="Log(In-State Tuition (USD))", title="(5b) Log(In-State Tuition):
Private vs Public")
ggsave(filename="writeup/5b_log.png", width=6.125, height=3.5, units="in")

# check group standard errors
tuitionData.b %>%
  group_by(Type) %>%
  summarize(sd=sd(LogInState))

```

```

# Perform a two-sample t-test
tt.b <- t.test(formula=LogInState~Type, data=tuitionData.b,
               var.equal=TRUE, conf.level=0.95, alternative="greater")
tt.b

# take antilog of the estimate
exp(-diff(tt.b$estimate)[[1]])

# Perform a two sided t-test for the confidence interval and take antilog
exp(t.test(formula=LogInState~Type, data=tuitionData.b,
           var.equal=TRUE, conf.level=0.95, alternative="two.sided")$conf.int)

# Part c: Out-of-state tuition (private > public) ### ### ### ### ### ### ###
### ### ### ### ### ### ### ### ### ### ### ### ###

# Create dataset for part c
tuitionData.c <- tuitionData %>%
  mutate(InState=NULL)

# create boxplots
ggplot(tuitionData.c, aes(x=Type, y=OutOfState)) +
  geom_violin(alpha=0.15) +
  geom_boxplot() +
  labs(y="Out-of-State Tuition (USD)", title="(5c) Out-of-State Tuition:
Private vs Public")
ggsave(filename="writeup/5c_original.png", width=6.125, height=3.5,
units="in")

# Create a log-transformed OutOfState tuition column
tuitionData.c <- tuitionData.c %>%
  mutate(LogOutOfState=log(OutOfState))

# create boxplots of the log-transformed data
ggplot(tuitionData.c, aes(x=Type, y=LogOutOfState)) +
  geom_violin(alpha=0.15) +
  geom_boxplot() +
  labs(y="Log(Out-ofn-State Tuition (USD))", title="(5c) Log(Out-of-State
Tuition): Private vs Public")
ggsave(filename="writeup/5c_log.png", width=6.125, height=3.5, units="in")

# check group standard errors
tuitionData.c %>%
  group_by(Type) %>%
  summarize(sd=sd(LogOutOfState))

# Perform a two-sample t-test
tt.c <- t.test(formula=LogOutOfState~Type, data=tuitionData.c,
               var.equal=TRUE, conf.level=0.95, alternative="greater")
tt.c

# take antilog of the estimate
exp(-diff(tt.c$estimate)[[1]])

# Perform a two sided t-test for the confidence interval and take antilog

```

```

exp(t.test(formula=LogOutOfState~Type, data=tuitionData.c,
           var.equal=TRUE, conf.level=0.95, alternative="two.sided")$conf.int)

# Check if outliers impace results in part c ## ## ## ## ## ## ## ## ## ## ## ## ## ##
## ## ## ##

# Create a dataset for part c excluding outliers
tuitionData.c2 <- tuitionData.c %>%
  filter(Type=="Private" & LogOutOfState < 11 |
         Type=="Public" & LogOutOfState < 10 & LogOutOfState > 9)

# Perform a two-sample t-test
tt.c2 <- t.test(formula=LogOutOfState~Type, data=tuitionData.c,
               var.equal=TRUE, conf.level=0.95, alternative="greater")
tt.c2

# take antilog of the estimate
exp(-diff(tt.c2$estimate)[[1]])

# Perform a two sided t-test for the confidence interval and take antilog
exp(t.test(formula=LogOutOfState~Type, data=tuitionData.c2,
           var.equal=TRUE, conf.level=0.95, alternative="two.sided")$conf.int)

# outliers don't impact the conclusion, so we only report the results that
included outliers

rm(list = ls()) # clear working environment


# Problem 6: Ramsey 4.19
#####

# load data
sparrowData <- Sleuth3::ex0221

# perform rank-sum test
wilcox.test(formula=Humerus~Status, data=sparrowData,
            paired=FALSE, # perform rank-sum test
            correct=TRUE, # apply continuity correction in the normal approx
            for the p-value
            conf.int=TRUE, conf.level=0.95)

# check if group sd are nearly equal
sparrowData %>%
  group_by(Status) %>%
  summarize(count=n(),
            mean=mean(Humerus),
            sd=sd(Humerus))

# check if group sd are nearly equal when excluding the outlier
sparrowData %>%
  filter(!(Humerus==min(Humerus) & Status=="Perished")) %>%
  group_by(Status) %>%
  summarize(count=n(),

```

```
mean=mean(Humerus),  
sd=sd(Humerus))  
  
rm(list = ls()) # clear working environment
```