

STAT W4201 001, Homework 4

Brian Weinstein (bmw2148)

Feb 24, 2016

Code is attached here and also posted at <https://github.com/BrianWeinstein/advanced-data-analysis>. Where relevant, code snippets and output are included in-line.

Problem 1: Ramsey 5.23

The data provides overwhelming evidence that the mean oxygen isotopic composition in the 12 bone samples are different (a p-value of 9.7×10^{-7} from a one-way analysis of variance (ANOVA) F-test).

The ANOVA table testing for a difference in mean oxygen isotopic composition is shown below, and a boxplot of oxygen composition for each bone is shown in Figure 1.

Source of Variation	Sum of Squares	d.f.	Mean Square	F-Statistic	p-Value
Between Groups	6.0675	11	0.55159	7.4268	9.73×10^{-7}
Within Groups	2.9708	40	0.07427		
Total	9.0383	51			




Figure 1: Oxygen Isotopic Composition (per mil deviations from SMOW) for twelve bones of a single Tyrannosaurus rex specimen.

Problem 2: Ramsey 5.25

- (a) *How strong is the evidence that at least one of the five population distributions (corresponding to the different years of education) is different from the others?*

Figure 2 shows the distribution of income for 5 different “years of education” groupings. The boxplots show (1) the presence of severe outliers, and (2) that the group standard deviations increase as the years of education increases.




Figure 2: Income vs Years of Education for 2,584 observations among 5 “Years of Education” groupings.

Examining the dataset, we also see that the group sample sizes are different.

```
> # check group sample sizes and standard deviations
> incomeEduData %>%
+   group_by(Educ) %>%
+   summarize(numObs=n(), mean=mean(Income2005),
+             median=median(Income2005), stdev=sd(Income2005))
Source: local data frame [5 x 5]
```

	Educ	numObs	mean	median	stdev
	(fctr)	(int)	(dbl)	(dbl)	(dbl)
1	<12	136	28301.45	23500	21021.90
2	12	1020	36864.90	31000	29369.73
3	13-15	648	44875.96	38000	33913.54
4	16	406	69996.97	56500	64256.80
5	>16	374	76855.46	60500	65428.29

The F-tests are not robust to a lack of equal standard deviations and are not resistant to severe outliers. Since the education groups with higher mean incomes also have higher spreads, this dataset is a good candidate for a log transformation.

On the log scale, there are fewer outliers (which are also less-severe) and the standard deviations are nearly equal, as shown in Figure 3 and in the table below.

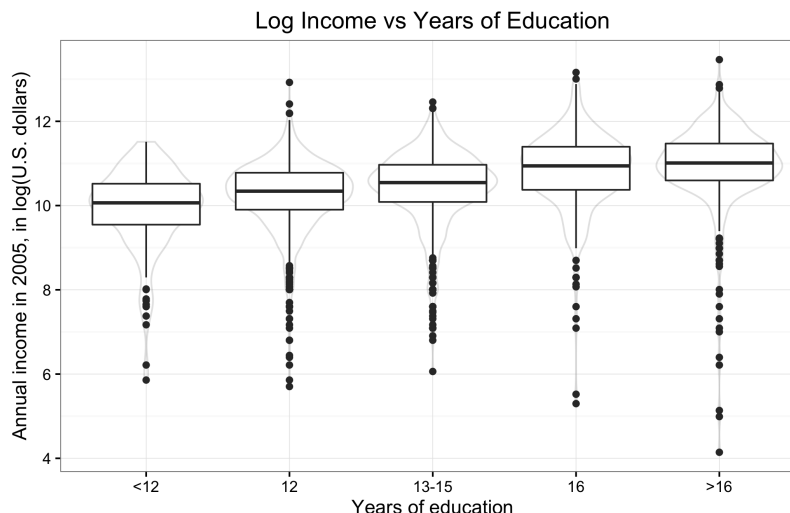


Figure 3: Log(Income) vs Years of Education for 2,584 observations among 5 “Years of Education” groupings.

```
> # check group sample sizes and standard deviations on log scale
> incomeEduData %>%
+   group_by(Educ) %>%
+   summarize(numObs=n(), mean=mean(LogIncome2005),
+             median=median(LogIncome2005), stdev=sd(LogIncome2005))
Source: local data frame [5 x 5]
```

	Educ (fctr)	numObs (int)	mean (dbl)	median (dbl)	stdev (dbl)
1	<12	136	9.89934	10.06453	0.9988809
2	12	1020	10.22721	10.34174	0.8539854
3	13-15	648	10.39121	10.54534	0.9288173
4	16	406	10.79709	10.94196	0.9581051
5	>16	374	10.89790	11.01036	1.0665910

Performing the one-way ANOVA F-test on the log-transformed incomes we find overwhelming evidence that at least one of the five population distributions is different from the others. The ANOVA table is shown below.

Source of Variation	Sum of Squares	d.f.	Mean Square	F-Statistic	p-Value
Between Groups	217.65	4	54.413	62.87	2.2×10^{-16}
Within Groups	2232.12	2579	0.865		
Total	2449.774	2583			

Although it isn't entirely justified here (as per Display 3.6), when performing the ANOVA F-test on the dataset excluding outliers¹ we still find overwhelming evidence that at least one distribution is different from the others, so the results are not included here.

¹Here, an outlier is defined as an observation more than 1.5 times the group interquartile range below the first quartile or above the third quartile.

- (b) *By how many dollars or by what percent does the mean or median for each of the last four categories exceed that of the next lowest category?*

The `CompareTwoEducGroups` function (see attached code for function definition) performs a two-sample t-test to test the hypothesis that the mean log income in the first specified “Years of Education” (YOE) group is greater than the mean log income in the second specified group. It outputs a one-sided p-value, an estimated value for the multiplicative treatment effect (in USD — the original scale), and a 95% confidence interval for the multiplicative treatment effect (also in USD).

- i. (**>16**) vs (**16**)

```
> CompareTwoEducGroups(data_frame=incomeEduData, Educ_groups=c(">16", "16"))
oneSidedPVal      estimate confInt_lower confInt_upper
0.08238005      1.10607335  0.95934487  1.27524345
> CompareTwoEducGroups(data_frame=incomeEduDataExclOutliers, Educ_groups=c(">16", "16"))
oneSidedPVal      estimate confInt_lower confInt_upper
3.968804e-05      1.224599e+00  1.107807e+00  1.353704e+00
```

The data provides little evidence that the >16 YOE population earns a higher income than the 16 YOE population (one-sided p-value 0.08238; two-sample t-test). A 95% confidence interval for the number of times by which the >16 YOE income exceeds the 16 YOE income is 0.95934 to 1.27524 times.

When excluding outliers, however, (see Figure 4) the data provides convincing evidence that the >16 YOE population earns a higher income than the 16 YOE population (one-sided p-value 3.97×10^{-5} ; two-sample t-test). Income is estimated to be 1.22460 times greater for the those with >16 YOE compared to those with 16 YOE, with a 95% confidence interval of 1.10781 to 1.35370 times (i.e., an estimated 22% increase; 95% CI from 11% to 35%).




Figure 4: $\text{Log}(\text{Income})$ vs Years of Education for 2,432 observations (i.e., excluding 152 outliers) among 5 “Years of Education” groupings.

- ii. (**16**) vs (**13–15**)

```
> CompareTwoEducGroups(data_frame=incomeEduData, Educ_groups=c("16", "13-15"))
oneSidedPVal      estimate confInt_lower confInt_upper
7.649007e-12  1.500615e+00  1.335230e+00  1.686486e+00
```

The data provides convincing evidence that the 16 YEO population earns a higher income than the 13–15 YOE population (one-sided p-value 7.65×10^{-12} ; two-sample t-test). Income is estimated to be 1.50062 times greater for the those with 16 YOE compared to those with 13–15 YOE, with a 95% confidence interval of 1.33523 to 1.68649 times (i.e., an estimated 50% increase; 95% CI from 34% to 69%).

iii. (13–15) vs (12)

```
> CompareTwoEducGroups(data_frame=incomeEduData, Educ_groups=c("13-15", "12"))
oneSidedPVal      estimate confInt_lower confInt_upper
0.0001140482  1.1782093370  1.0799490482  1.2854099405
```

The data provides convincing evidence that the 13–15 YEO population earns a higher income than the 12 YOE population (one-sided p-value 0.00011; two-sample t-test). Income is estimated to be 1.17821 times greater for the those with 13–15 YOE compared to those with 12 YOE, with a 95% confidence interval of 1.07995 to 1.28541 times (i.e., an estimated 17% increase; 95% CI from 8.0% to 29%).

iv. (12) vs (<12)

```
> CompareTwoEducGroups(data_frame=incomeEduData, Educ_groups=c("12", "<12"))
oneSidedPVal      estimate confInt_lower confInt_upper
0.0000204658  1.3880147741  1.1872749406  1.6226949185
```

The data provides convincing evidence that the 12 YEO population earns a higher income than the <12 YOE population (one-sided p-value 0.00002; two-sample t-test). Income is estimated to be 1.38801 times greater for the those with 12 YOE compared to those with <12 YOE, with a 95% confidence interval of 1.18727 to 1.62269 times (i.e., an estimated 39% increase; 95% CI from 19% to 62%).

Problem 3: Ramsey 6.12

The quantity of interest is the linear combination

$$\gamma = \frac{\mu_{\text{amputee}} + \mu_{\text{crutches}} + \mu_{\text{wheelchair}}}{3} - \frac{\mu_{\text{hearing}}}{1},$$

which we estimate by

$$g = \frac{\bar{Y}_{\text{amputee}} + \bar{Y}_{\text{crutches}} + \bar{Y}_{\text{wheelchair}}}{3} - \frac{\bar{Y}_{\text{hearing}}}{1},$$

where the amputee, crutches, and wheelchair groups are handicaps of “mobility”, and the hearing group is a handicap of “communication”. Since the coefficients sum to 0, g is a linear contrast.

We’re testing whether the average of the mean scores for the mobility group is equal to the mean score of the communication group. That is, $H_0 : g = 0$, $H_A : g \neq 0$.

Using the `gmodels::fit.contrast` R function, the data provides moderate evidence that the average of the mean scores for the mobility group is not equal to the mean score of the communication group (two-sided p-value 0.02217 for a linear contrast). The score is estimated to be 1.18095 points higher for the mobility group compared to the communication group (95% confidence interval 0.17452 to 2.18738 points).

```

> # check the order of Handicap factor levels
> levels(handicapData$Handicap)
[1] "Amputee" "Crutches" "Hearing" "None" "Wheelchair"
>
> # test if the the avg of score means for
> # amputee/crutches/wheelchair is equal to to hearing
> fit.contrast(model=lm(Score ~ Handicap, data=handicapData),
+             varname="Handicap", coeff=c(1/3, 1/3, -1, 0, 1/3),
+             conf.int=0.95, df=TRUE)

```

		Estimate	Std. Error	t value	Pr(> t)	DF	lower CI	upper CI
Handicap	c=(0.333 0.333 -1 0 0.333)	1.180952	0.5039353	2.34346	0.02217534	65	0.174524	2.187381

Problem 4: Ramsey 6.15

Test scores for the experimental CAD instruction course are shown below:

Group	Logo	Teaching method	<i>n</i>	Average	SD
1	L + D	Lecture and discussion	9	30.20	3.82
2	R	Programmed text	9	28.80	5.26
3	R + L	Programmed text with lectures	9	26.20	4.66
4	C	Computer instruction	9	31.10	4.91
5	C + L	Computer instruction with lectures	9	30.20	3.53

- (a) Compute the pooled estimate of the standard deviation from these summary statistics.

The pooled estimate of the standard deviation s_p is given by

$$\begin{aligned}
 s_p &= \sqrt{\frac{\sum_{i=1}^5 (n_i - 1) s_i^2}{\sum_{i=1}^5 (n_i - 1)}} \\
 &= \sqrt{\frac{(9 - 1)(3.82)^2 + (9 - 1)(5.26)^2 + (9 - 1)(4.66)^2 + (9 - 1)(4.91)^2 + (9 - 1)(3.53)^2}{(9 - 1) \cdot 5}} \\
 &= 4.484297,
 \end{aligned}$$

and has d.f. = $9 + 9 + 9 + 9 + 9 - 5 = 40$ degrees of freedom.

- (b) Determine a set of coefficients that will contrast the methods using programmed text as part of the method (groups 2 and 3) with those that do not use programmed text (1, 4, and 5).

For groups 1 through 5, respectively, the coefficients C_i contrasting groups 2 and 3 with groups 1, 4, and 5 are

$$\left(\frac{1}{3}, -\frac{1}{2}, -\frac{1}{2}, \frac{1}{3}, \frac{1}{3}\right).$$

- (c) Estimate the contrast in (b) and compute a 95% confidence interval.

An estimate of the linear contrast of interest is

$$\begin{aligned}
 g &= \frac{\bar{Y}_{L+D} + \bar{Y}_C + \bar{Y}_{C+L}}{3} - \frac{\bar{Y}_R + \bar{Y}_{R+L}}{2} \\
 &= \frac{30.20 + 31.10 + 30.20}{3} - \frac{28.80 + 26.20}{2} \\
 &= 3.
 \end{aligned}$$

The standard error of the estimate is

$$\begin{aligned} \text{SE}(g) &= s_p \sqrt{\sum_{i=1}^5 \frac{C_i^2}{n_i}} \\ &= (4.484297) \sqrt{\frac{(1/3)^2}{9} + \frac{(-1/2)^2}{9} + \frac{(-1/2)^2}{9} + \frac{(1/3)^2}{9} + \frac{(1/3)^2}{9}} \\ &= 1.364528. \end{aligned}$$

A 95% confidence interval is

$$\begin{aligned} &g \pm t_{40}(0.975) \cdot \text{SE}(g) \\ &3 \pm 2.021075 \cdot 1.364528 \\ &3 \pm 2.757814 \\ &\Rightarrow 0.2421858 \leq g \leq 5.757814. \end{aligned}$$

Problem 5: Ramsey 6.16

A study involving 36 subjects randomly assigned six each to six treatment groups gives an ANOVA F-test with $p\text{-value} = 0.0850$. What multipliers are used to construct 95% confidence intervals for treatment differences with the following methods:

(a) **LSD**

For the LSD, the multiplier M in the interval half-width is given by the $100(1 - \frac{\alpha}{2})\%$ critical value in the t-distribution with degrees of freedom equal to those associated with the pooled SD (here, $\text{d.f.}_{s_p} = (6 + 6 + 6 + 6 + 6 + 6) - 6 = 30$).

Therefore, $M = t_{30}(1 - \frac{0.05}{2}) = t_{30}(0.975) = 2.042272$.

```
> # multiplier for LSD
> qt(p=(1-(0.05/2)), df=30)
[1] 2.042272
```

(b) **F-protected LSD**

Since the p-value from the F-test is large (>0.05), we can't declare any individual difference significant, and neither the Tukey-Kramer nor the Scheffe multipliers control either the individual or familywise success rate at 95%.

(c) **TukeyKramer**

In the Tukey-Kramer procedure, the multiplier M is given by $[q_{I,(n-I)}(1 - \alpha)]/\sqrt{2}$, where $[q_{I,(n-I)}(1 - \alpha)]$ is the quantile function of the studentized range distribution with I groups and $(n - I)$ degrees of freedom.

Therefore, $M = [q_{6,(36-I6)}(1 - 0.05)]/\sqrt{2} = 3.041594$.

```
> # multiplier for Tukey-Kramer
> qtukey(p=(1-0.05), nmeans=6, df=30) / sqrt(2)
[1] 3.041594
```

(d) **Bonferroni**

In the Bonferroni procedure, the multiplier M is given by $t_{\text{d.f.}}(1 - \frac{\alpha}{2k})$, where $k = \frac{I(I-1)}{2}$ is the number of pairs of means to be compared and d.f. is the degrees of freedom associated with the pooled SD.

Here, $k = \frac{6(6-1)}{2} = 15$, therefore, $M = t_{30}(1 - \frac{0.05}{2 \cdot 15}) = t_{30}(0.9983333) = 3.188806$.

```
> # multiplier for Bonferroni
> qt(p=(1-(0.05/(2*15))), df=30)
[1] 3.188806
```

(e) **Scheffe**

In the Scheffe procedure, the multiplier M is given by $\sqrt{(I-1)F_{(I-1),d.f.}(1-\alpha)}$, where $F_{(I-1),d.f.}(1-\alpha)$ is the $(1-\alpha)$ th percentile of the F-distribution with $(I-1)$ and d.f. degrees of freedom.

Therefore $M = \sqrt{(6-1)F_{(6-1),30}(1-0.05)} = \sqrt{5 \cdot F_{5,30}(0.95)} = \sqrt{5 \cdot 2.533555} = 3.559181$.

```
> # multiplier for Scheffe
> sqrt(5 * qf(p=(1-0.05), df1=5, df2=30))
[1] 3.559181
```

Problem 6: Ramsey 6.23

The data provides moderately strong, but not convincing evidence that there are differences in average weight loss among these diets (a p-value of 0.04086 from a one-way ANOVA F-test).

The ANOVA table testing for a difference in weight loss after 24 months of dieting is shown below, and a boxplot of these weight losses is shown in Figure 5.

Source of Variation	Sum of Squares	d.f.	Mean Square	F-Statistic	p-Value
Between Groups	216.9	2	108.430	3.2358	0.04086
Within Groups	9013.9	269	33.509		
Total	9230.753	271			

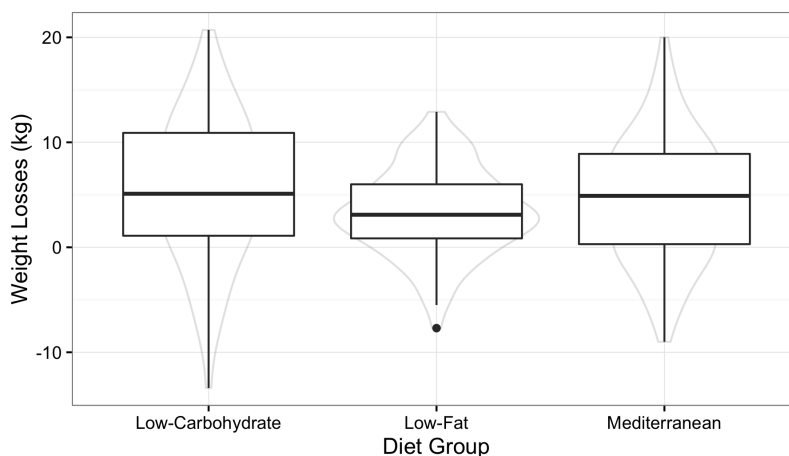


Figure 5: Weight losses (in kg) of 272 participants who dieted for two years on one of three different diets.

The difference between the average weight loss in the Low-Carbohydrate group and the Low-Fat group is difficult to attribute to chance. The weight loss in the Low-Carbohydrate group

is estimated to be 2.1828 kg higher than in the Low-Fat group, with a 95% confidence interval from 0.1408 to 4.2248 kg based on the Tukey-Kramer procedure.

The difference between the average weight loss in the Mediterranean and Low-Fat groups, and the difference between the average weight loss in the Low-Carbohydrate and Mediterranean groups are not significant.

```
> TukeyHSD(x=aov(lm(WtLoss24 ~ Group, data = dietData)), which="Group",
+ ordered=TRUE, conf.level=0.95)
  Tukey multiple comparisons of means
    95% family-wise confidence level
    factor levels have been ordered
```

```
Fit: aov(formula = lm(WtLoss24 ~ Group, data = dietData))
```

\$Group	diff	lwr	upr	p adj
Mediterranean-Low-Fat	1.2978952	-0.6974180	3.293208	0.2771180
Low-Carbohydrate-Low-Fat	2.1828035	0.1408361	4.224771	0.0329364
Low-Carbohydrate-Mediterranean	0.8849083	-1.1622656	2.932082	0.5656813